


Operating Systems (234123)

Virtual memory

Dan Tsafir (19/5/2014, 26/5/2014, 2/6/2014)
Partially based on slides by Hagit Attiya

DRAM (dynamic random-access memory)


- This is the physical memory...
 - Patriot Signature Apple DDR3 16GB (2 X 8GB) 1600MHz PC3-12800



DRAM (dynamic random-access memory)


- **This is the physical memory...**

- Patriot Signature Apple DDR3 16GB (2 X 8GB) 1600MHz PC3-12800
- Price
 - May 2013: \$134 (~ \$8 per 1 GB)
 - May 2014: \$159 (~ \$10 per 1GB) ☺
- Bandwidth
 - 1600 (MHz; transfers per sec)
 - x 64 (bit per bus transfer, for DDR3)
 - / 8 (bits per byte)
 - = $1600 \times 10^6 \times 64 / 8$ [B/sec]
 - = 12,800 MB/sec = 12.8 [GB/sec]
 - (= meaning of the “PC3-12800”; GB= 10^9 , in this case)
- Latency
 - A handful of 100s of CPU cycles (\approx 100s of ns)
 - (Assuming not cached)



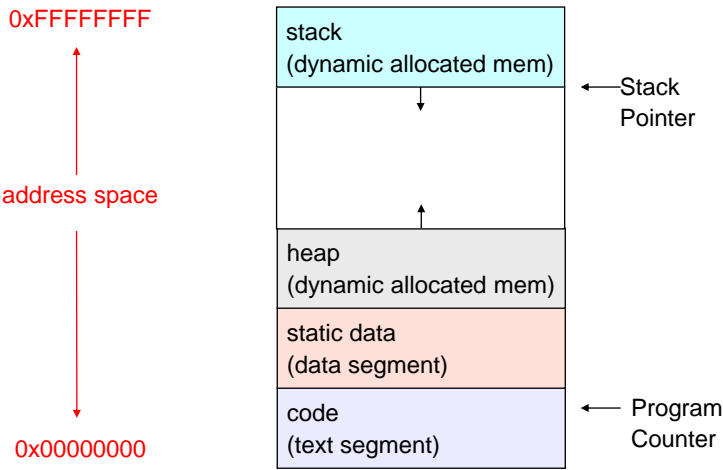
Reminder: in comparison to HDD...

	HDD (hard disk drive)	DRAM	times faster / better
availability	non-volatile	volatile	
sequential access bandwidth	≈ 100 MB/s	≈ 10 GB/s	$> \times 100$
random access bandwidth (reading 4KB chunks at a time)	≈ 1 MB/s	≈ 10 GB/s	$> \times 10,000$
latency	≈ 5 ms	≈ 100 ns	$> \times 50,000$



Reminder

- This is how a program believes its memory looks like...



But how can this be?

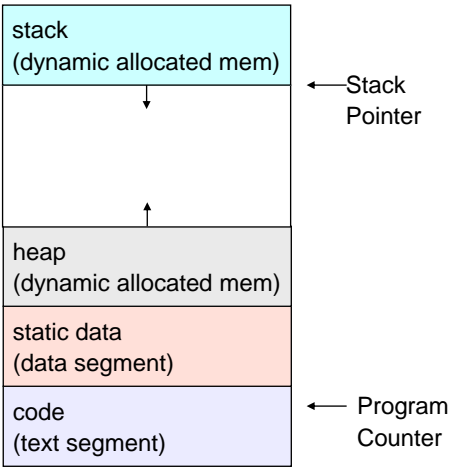
- It's impossible for all processes to share these addresses

⇒ They translate to different locations, so...

⇒ **There's an abstraction layer!**

⇒ This abstraction is called "virtual memory"

- Provides this —————→ illusion to all programs
- And more. Specifically... [next page]



Virtual memory (vmem) – motivation

- **Per-program Illusion of contiguous memory**
 - Programmers need not worry about *where* data is placed exactly
 - They use the ideal address space from 2 slides ago
- **Isolation between processes**
 - Processes can concurrently run on the same processor
 - Yet vmem prevents them from accessing the memory of one another
 - (But still allows for convenient sharing when required)
- **Dynamic growth**
 - Can add memory to process's heap/stack at runtime, as needed
- **Illusion of large memory => memory overcommitment**
 - DRAM often: (i) one of the most costly parts; (ii) the bottleneck resource
 - Vmem size can be bigger than physical memory size
 - Sum of vmem spaces (across all processes) can be \geq physical
- **Access control**
 - Decide if individual memory chunks can be read / written / executed

OS (234123) - virtual memory

7


HOW VMEM WORKS, IN PRINCIPLE

OS (234123) - virtual memory

8

Virtual memory – terminology

- **Virtual address (VA)**
 - Used by the program/programmer
 - “Ideal” = contiguous & as big as we’d like
- **Physical address (PA)**
 - The real, underlying physical memory address
 - Completely abstracted away by OS/HW
- **Memory (virtual & physical) is divided into fixed size blocks**
 - “Page” = chunk of contiguous data (in virtual or physical space)
 - “Frame” = physical memory exactly big enough to hold one page
 - $|page| = |frame| = 2^k$ (bytes)
 - Typically, $k = 12$, namely a page (and frame) size is 4KB
- **Pages & frames are always aligned on 4KB boundaries**
 - Both in physical and virtual memory spaces: 0, 4KB, 8KB, 12KB, 16KB, ...




OS (234123) - virtual memory 9

Virtual memory – basic idea

- “Map” pages to frames, such that VA space is contiguous
 - Pages can be “mapped” into (associated with) arbitrary frames at arbitrary locations
- Pages can reside in memory, or on disk
 - Hence, we achieved the aforementioned memory overcommitment
- All programs are written using VAs
 - And somehow VAs are seamlessly translated into PAs
 - As we will see later, **translation is a HW/SW mechanism**, whereby
 - OS sets the VA=>PA mappings
 - HW does on-the-fly translation from VA to PA

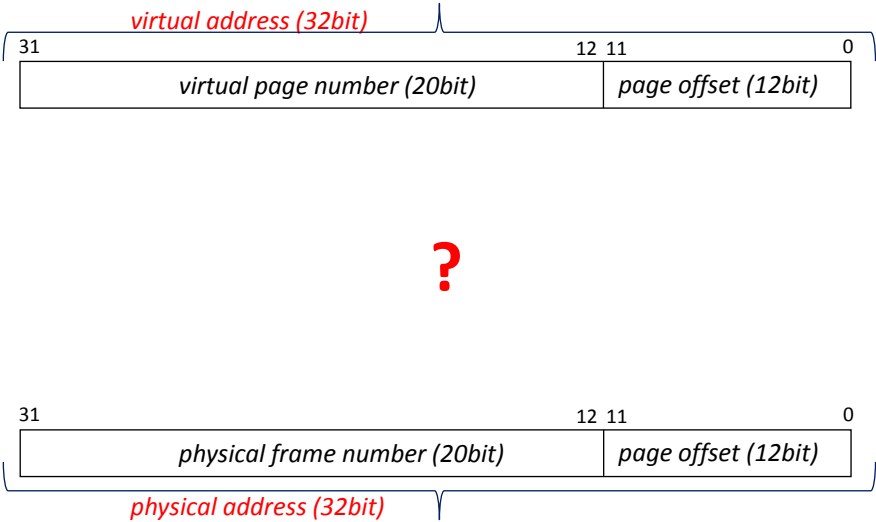
OS (234123) - virtual memory 10

Per-process virtual memory simplistic illustration

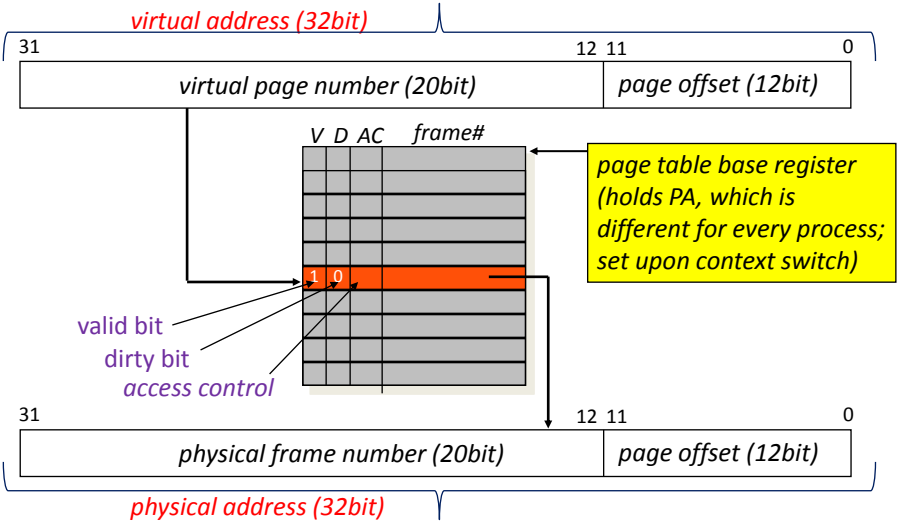


- Memory may serve as a “cache” for secondary storage (disk);
- Immediate advantages
 1. Illusion of contiguity & of having more physical memory
 2. Program’s actual location unimportant
 3. Dynamic growth, isolation, & sharing are easy to obtain


How to map (assume 32bit address)?



Use per-process “page table” (in DRAM)



Per-process page table



Upon (each & every) memory access

- **If (valid == 1)**
 - Page is in main memory @ PA stored in table
 - Data is readily available & can be used
- **else**
 - Suspend process
 - Fetch page from disk
 - Resume process, which will re-execute faulting instruction
 - Now it'll succeed
- **Access Control**
 - R=read-only, R/W=read/write, X=execute
 - If access type incompatible with specified access rights
 - ⇒ protection violation fault ⇒ interrupt ⇒ signal

OS (234123) - virtual memory

15

Major page fault

- **Page not in memory ⇒ need to retrieve it from disk**
 1. CPU detects the situation (valid bit = 0)
 - But it cannot remedy the situation on its own
 - CPU doesn't communicate with disks; OS does that
 2. CPU generates interrupt and transfers control to the OS
 - Invoking the OS page-fault handler
 3. OS regains control & initiates I/O read operations
 - To read missing page from disk to DRAM
 - Possibly need to write victim page to disk (if no room & dirty)
 4. OS suspends process & context switches to another
 - It will take milliseconds for I/O operations to complete
 5. Upon read completion, OS makes suspended process runnable again
 - It'll soon be chosen for execution
 6. When process is resumed, faulting operation is re-executed
 - Now it will succeed because the page is there

OS (234123) - virtual memory

16

Minor page fault

- **Not all page faults are major...** end of 1st lecture
 - That is, not all page faults take milliseconds to handle
- **Sometimes the info accessed by the program is in DRAM**
 1. Yet the 'valid' bit is off, or
 2. The 'valid' bit is on, but there's an access violation
- **For example**
 1. A memory page is found in the "buffer cache" of the OS
 - Caching previously read files, possibly by other processes
 - All read/write ops of all processes go through the buffer cache
 2. **COW** (copy-on-write)
 - Used for, e.g., implementing fork()
 - In child, map all pages to parent memory space, but for reading
 - When child writes => page fault => create a private copy for it
- **Minor page faults are much faster to handle**
 - Processes usually not suspended due to minor page faults

OS (234123) - virtual memory

17

Related concepts

- **"Page in" & "page out" a chunk of data**
 - Page in ⇔ copy page from disk to DRAM (= read)
 - Page out ⇔ copy page from DRAM to disk (=write)
- **The mmap() system call**
 - (Mmap = memory map)
 - Map a given file into the virtual memory space of the process
 - A file becomes an "array of bytes" (backed by disk), and gets a VA
 - With the right mmap() flags, reading/writing from/to the array translates to reading/writing from/to the file
 - A mmap()ed file is said to be a "memory-mapped file"

OS (234123) - virtual memory

18

Related concepts

- **Anonymous pages**
 - Heap/stack pages; not file
 - Can be allocated by mmap (MAP_ANONYMOUS)
 - Namely, can implement malloc with mmap, which can be invoked whenever malloc runs out of memory
 - Typically, however, malloc is implemented using the system calls brk() and sbrk()
- **Named pages**
 - Backed by a file (via mmap)
- **The mmap system call can also map file pages “privately”**
 - Using the MAP_PRIVATE flag
 - Meaning, if the page is changed, this doesn’t affect the underlying file
 - There will be a copy-on-write

OS (234123) - virtual memory

19

Related concepts

- **Q/A**
 - Q: do we need to page out named pages?
A: only if they are “dirty”
 - Q: when we read() a file page (with the read syscall), is it named?
A: no, as changes to the memory buffer will not affect the file;
need to write() to the file in order to make a change
 - Q: can anonymous pages reside on disk?
A: yes [see next slide]

OS (234123) - virtual memory

20

Related concepts

- **Swap space**
 - Disk area (file) where anonymous pages are written, if the OS decides they have no room in DRAM
 - Page is said to be “swapped out” when this occurs (and “swapped in” for the reverse operation)
 - Swap area contains anonymous pages (including mmap anonymity)
- **Swapping vs. paging**
 - In the olden days, “swapping in/out” referred to the entire memory of a process (not just to a certain page)
 - Nowadays people typically use them interchangeably
 - <http://en.wikipedia.org/wiki/Paging> makes the following distinction
 - Page in = transfer page from anywhere on disk to DRAM
 - Swap in = transfer a page from swap-space to DRAM
 - But we do not use this distinction

OS (234123) - virtual memory

21

Related concepts

- **Demand-paging**
 - OS reads a page from disk into DRAM only if the process attempts to access it (and, hence, a page fault occurs)
 - That is, OS pages data in only via page faults (+ prefetching)
 - It follows that a process begins execution with most of its pages not residing in physical memory, and page faults occur until most of the data it uses is located in DRAM
 - Also called “lazy” loading
 - Q: what’s the benefit?
A: reading is costly, & with demand-paging we only read what we need
- **Readahead prefetching (anticipatory paging)**
 - read() does prefetching when identifying sequential access
 - Page fault handler does the same
 - Complements demand-paging in an attempt to minimize page faults

OS (234123) - virtual memory

22

Related concepts

- **Working set**

- $WS_p(k)$ = Pages accessed by P in the last k accesses

time	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
page	2	6	1	5	7	7	7	7	5	1	6	2	3	4	4	4	3
$WS(k=3)$	2	6	1	5	7	7	7	7	5	1	6	2	3	4	4	4	3
		2	6	1	5	5			7	5	1	6	2	3	3		4
			2	6	1				7	5	1	6	2				

- For a fixed value of k, smaller $WS_p(k)$ indicates more locality
- If the current working set is not found in memory then...

- **Thrashing**

- When we've overcommitted too much memory and there isn't enough physical memory, the system might enter a state of thrashing, that is
- Virtual memory is in a constant state of paging, rapidly exchanging data between memory & disk
- Nearly nothing else is done in the system (causes performance to degrade or collapse)

OS (234123) - virtual memory

23

Did we achieve our goals?

- **Illusion of contiguous memory**
 - Yes: virtual memory is contiguous by definition
- **Illusion of large memory (possibly bigger than physical mem)**
 - Yes: chunks that don't fit into physical memory reside on disk
- **Dynamic growth**
 - Yes: heap & stack can grow at runtime by mapping more VAs to PAs, in order
- **Isolation between processes**
 - Yes: same VAs point to different PAs; as long as we keep PAs disjoint on a per-process basis, the processes are isolated
- **Memory overcommitment**
 - Yes: using the disk, $\sum_{i=1}^n vmem_i$ (for n processes) can be > physical size
- **Access control**
 - Yes: HW enforces PTE (= page table entry) bits; e.g., it will reject a write to a page that is marked 'read-only'

OS (234123) - virtual memory

24

But how does it perform?

- **(1) Temporal locality helps (“temporal” relates to time)**
 - Typically, during a given time interval, a process uses only a fraction of its memory, over and over again
 - So it’s fine to keep currently unused parts on disk
 - As long as the working set is mostly in memory
- **(2) Asynchronous I/O helps**
 - Writes are non-blocking: when writing a page to disk, we don’t need to block the associated process
 - When reading stuff, we can run other processes
- **(3) Demand-paging helps**
 - Pages fetched from secondary memory only upon the first page fault, rather than, e.g., upon file open (we bring only what we need)

OS (234123) - virtual memory

25

But how does it perform?

- **(4) Special locality helps (“special” relates to space)**
 - When reading page P from disk, we typically employ the **readahead** optimization, namely, we also bring some additional file pages that come immediately after P, even though they weren’t accessed yet
 - As we’ve learned, “**locality principle**” suggests these pages would be used soon
 - Locality = special locality + temporal locality
 - Locality principle contends that most programs exhibit special and temporal locality when utilizing the memory
 - So some pages don’t induce a page fault when accessed for the first time and are simply there

OS (234123) - virtual memory

26

But how does it perform?


- **(5) Making VA=>PA translations fast helps**
 - The TLB (translation lookaside buffer) is a very small, very fast HW structure that caches recently used VA=>PA mappings
 - (e.g., 32-64-128 entries)
 - Given a VA, HW first searches for its translation in the TLB; and only if it's not there HW access the in-memory page table
 - Accessing the TLB takes very little time (e.g., a cycle)
 - Even though the TLB is very small, locality principle typically ensures it is rather effective
 - Special locality: since we work in 4KB page granularity, lots of nearby accesses fall in the same page
 - Temporal locality: same pages are used repeatedly
 - Lots of workloads (though certainly not all) approach 100% TLB hit rate

OS (234123) - virtual memory

27

VA => PA translation with TLB

- **Page table resides in memory**
 - Each translation requires a memory access
 - Might be required for each load/store!
- **TLB**
 - Cache recently used PTEs => speeds up translation
 - TLB access time comparable to L1
 - Typically:
 - HW fills TLB automatically by reading the page table on its own (no SW involvement)
 - OS can invalidate TLB entries (when should it?)
 - Processes are completely unaware of TLB



OS (234123) - virtual memory

28

But how does it perform?

- **(6) Using an intelligent page replacement policy helps**
 - When we need to evict a page from memory to disk
 - Page replacement policy decides which page it'll be
 - Also called “page reclamation”
 - Goal
 - Minimize number of future page faults
 - Minimize price of paging (evicting dirty pages costs more)
 - Typically done via a daemon process (“swapper”) that runs in the background
 - Start: when number of free frames drops below some threshold
 - Stop: when number of free frames exceeds some threshold

OS (234123) - virtual memory

29

Page reclamation algorithms

- Belady (optimal off-line; if we know the future; theoretical)
 - Greedily page out page accessed furthest in the future
- Could be FIFO (first-in first-out)
 - Simplest (no need to update upon each mem ref), but ignores usage
- LRU (least recently used) is better
 - But typically too wasteful (updated upon each mem ref)
- Second-chance
 - Set per-page “was it referenced?” bit
 - Can be done by HW or SW (how?)
 - Page out pages with bit=0 only, FIFO order
 - When traversed, if bit=1, set it to 0, and push the associated page to end of the list (in FIFO terms, page becomes newest)
- Clock
 - More efficient/popular variant of second-chance
 - Pages are cyclically ordered (no FIFO); search clockwise for first page with bit=0; set bit=0 for pages that have bit=1; periodically turn off 1s

OS (234123) - virtual memory

30

Page reclamation algorithms

- **NRU (not recently used)**
 - More sophisticated LRU approximation
 - HW or SW maintains per-page 'referenced' & 'modified' bits
 - Periodically (a few clock interrupts), SW turns 'referenced' off
 - Though if the modified bit is on, it stays on
 - Replacement algorithm partitions pages to
 - Class 0: not referenced, not modified
 - Class 1: not referenced, modified
 - Class 2: referenced, not modified
 - Class 3: referenced, modified
 - Choose at random a page from the lowest class for removal
 - Underlying principles (order is important):
 - Prefer keeping referenced over unreferenced
 - Prefer keeping modified over unmodified

OS (234123) - virtual memory

31

Hierarchical translation (radix tree)

HOW IT WORKS IN X86

OS (234123) - virtual memory

32

32bit x86 address translation


- **32bit address means $2^{32} = 4\text{GB}$ address space**
 - There are $2^{20} = 1\text{MB}$ pages (1MB pages x 4KB per-page = 4GB)
- **The job of the x86 virtual memory subsystem**
 - Translate 20 bits (= virtual page #) to 20 bits (= physical frame #)
 - Pages are 4KB-aligned, so it's enough to identify them with 20bits
- **Every process has a “page directory” (4KB page)**
 - Holds 1024 PDEs (page-directory entries)
 - Each PDE is comprised of 4 bytes (= 32 bits; $1024 \times 4 = 4\text{KB}$)
 - Each PDE contains
 - Bits: **Present? Accessed? Dirty? Writable? User? Cache Disabled?** ...
 - 20bit physical page frame number:
 - Where to find the corresponding “page table”
- **Every PDE points to a 4KB “page table”**
 - Holds 1024 PTEs (page-table entries) with same bits + 20bit frame#
 - Frame# points to a program's memory page

OS (234123) - virtual memory

33

2-level hierarchy


- **CR3 register points to page-directory**
 - Physical address set on context switch
 - Per process (but threads share it)
- **DIR (10 bits)**
 - Index of PDE (4bytes) in page-directory array (there are 1024 PDEs)
 - Each PDE holds 20bit of 4KB-aligned physical frame# of a 4KB page table
- **TABLE (10 bits)**
 - Index of PTE (4bytes) in page-table array (there are 1024 PTEs)
 - Each PTE holds 20 bit of 4KB-aligned physical frame# of a 4KB “regular” page
- **OFFSET (12 bits)**
 - Offset within the selected 4KB page



OS (234123) - virtual memory

34

4KB-page PTE format



OS (234123) - virtual memory

35

Combining user/supervisor & global page bits

- **Problem: when a system call happens...**
 - OS must run and use its own internal data structures
 - But we don't want every system call to induce a context switch
- **Solution: map OS data structures to address space of all processes**
 - Use the "user/supervisor bit" to indicate that only the OS ("ring 0") can access this memory area, whereas user code ("ring 3") cannot
 - Further, set the "global page" bit in the PTEs associated with the OS memory, which would then leave the corresponding TLB translations valid across context switches

OS (234123) - virtual memory

36

Why hierarchical?

- **The alternative: one linear page table**
 - Requires 2^{20} (PTEs) \times 4 (bytes per PTE) = 4MB
 - Could be wasteful (recall, it's a per-process overhead)
 - And of course it'll be much, much worse (\Leftrightarrow impossible) for 64bit addresses (discussed later)
- **The hierarchical translation allows us to avoid wasting memory**
 - Page tables (2nd-level) are allocated on-demand only
 - Far less “internal fragmentation” when memory space is sparse
 - Internal fragmentation
= when parts of an allocation unit remain unused
 - (Is there no internal fragmentation when using hierarchical translation?)

OS (234123) - virtual memory

37

HW/OS cooperate

- **HW defines data structures**
 - Structure of hierarchy, PTE bits, etc.
- **OS determines most of the content of page directories & page tables**
 - It explicitly sets the values of the PDEs and PTEs
- **HW does the “table walk” automatically**
 - If VA \Rightarrow PA translation not found in the TLB
 - Called “TLB miss”
 - HW knows where to find the page directory (using CR3)
 - HW walks the tables, hierarchically, until it reaches the data page
 - It inserts the VA \Rightarrow PA translation to the TLB
 - When page-faulting operation resumes, there will be a TLB hit
- **HW also responsible for setting bits**
 - Accessed & dirty bits (Why? Recall that OS emulate this behavior)
 - OS is responsible for turning these bits off

OS (234123) - virtual memory

38

HW/OS cooperate

- **HW populates TLB**
 - Whenever TLB miss occurs
- **OS invalidates TLB entries**
 - E.g., upon context switch (if HW doesn't support PID in TLB; nowadays it typically does, but the number of PIDs is limited)
- **OS is also responsible for synchronizing between TLBs**
 - Each core has its own TLB
 - **TLB shutdown:** when, e.g., one core invalidates PTE or changes access, OS must sync TLBs on other cores

end of 2nd lecture
(1st bullet)

OS (234123) - virtual memory

39

Locality helps address translation too

- **Recall CPU reads data from DRAM in cache-line resolution**
 - Cache-line size = 32 bytes in x86 32bit
 - So whenever the CPU reads one PTE
 - It actually inserts another 7 PTEs to the cache (to L2, but not to L1)
 - $32\text{bytes} / 4\text{bytes} = 8$ PTEs
 - The eight PTEs reflect contiguous virtual space (special locality)
 - If/when CPU encounter a TLB miss on one of the 7, it'll find them in L2 and will not have to read it from DRAM
 - For x86 64 bit, PTE is twice as long (8 bytes), but cache line is twice as long too (64 bytes)

OS (234123) - virtual memory

40

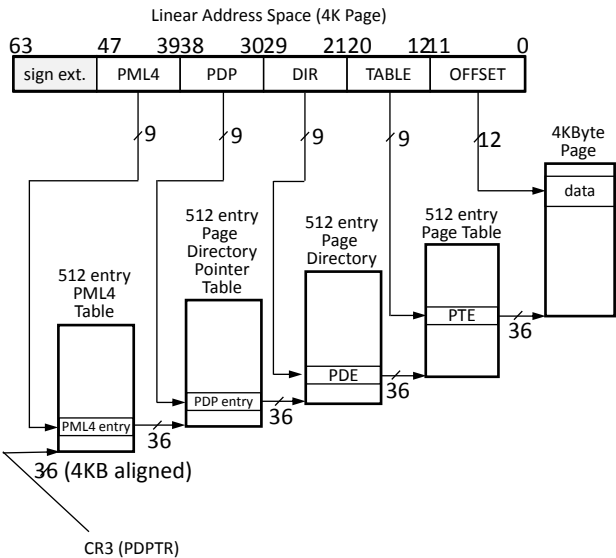
64bit x86 address translation

- 64bit address means $2^{64} = 16$ Exabyte address space**
 - $2^{20} \approx 10^6$ = Megabyte
 - $2^{30} \approx 10^9$ = Gigabyte
 - $2^{40} \approx 10^{12}$ = Terabyte
 - $2^{50} \approx 10^{15}$ = Petabyte
 - $2^{60} \approx 10^{18}$ = Exabyte (= a billion GBs)
 - ...DRAM can't be that big; 64 bits are thus much more than needed
- In practice, current x86_64 HW uses "only" 48 bits => 256 TB**
 - Still more than enough
- 48bit address reflects a 4-level hierarchy, divides into 5 parts**
 - 9bits X 4 levels + 12bits offset
 - Each PTE is 8 bytes (rather than 4, to be able to hold the wider address)
- The job of the x86 64bit virtual memory subsystem**
 - Translate 36 bits (= virtual page #) to 36 bits (= physical frame #)
 - As before, pages are 4KB-aligned

OS (234123) - virtual memory

41

4-level hierarchy



OS (234123) - virtual memory

42

More details:

<http://webcourse.cs.technion.ac.il/234123/Spring2013/ho/WCFiles/ppc64vm-zhang2009.pdf>

<http://webcourse.cs.technion.ac.il/234123/Spring2013/ho/WCFiles/ppc64vm-peng1995.pdf>

A different HW/SW virtual memory implementation

64-BIT POWER-PC (PPC) VIRTUAL MEMORY

OS (234123) - virtual memory

43

HW-support for virtual memory can be....

- **Implemented completely differently**
 - Not carved in stone
 - It's the HW vendor that decides
- **Intel defined in one way for the x86 architecture**
 - Radix tree page table
- **IBM defined it differently for the POWER architecture**
 - Hashed page table
 - Which we're going to describe next...

OS (234123) - virtual memory


44

3 address types: effective => virtual => real


- **Effective**
 - Each process uses 64-bit “effective” addresses
 - Effective addresses aren’t unique per-process
 - More or less equivalent to x86 “virtual” addresses
 - Get translated to PPC “virtual” addresses
- **Virtual**
 - A huge 80-bit address space
 - All processes live in and share this (single) space
 - Namely, if two processes have a page with the same virtual address
 - Then it’s the same page (= a shared page)
 - Not equivalent to x86 virtual addresses
 - Get translated into physical (“real”) address
- **Physical (a.k.a. real)**
 - 62-bit

PPC Segments

- **Effective & virtual spaces are partitioned into contiguous segments of 256MB (= 2^28 = 2^16*2^12 = 64K*4KB pages)**
 - Each segment is contiguous in the per-process effective memory space
 - Each segment is contiguous in the single, huge virtual space
 - Segments are 256MB-aligned, and can be private or shared (why?)
- **How many segments can there be in effective space?**
 - Effective space size is 2^64; so can be 2^(64-28) = 2^36 segments

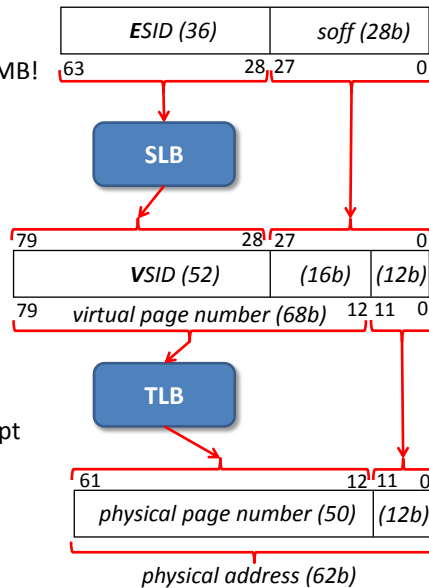


- **How many segments can there be in virtual space?**
 - Effective space size is 2^80; so can be 2^(80-28) = 2^52 segments



PPC SLB (segment lookaside buffer)

- **SLB is a fast, small HW cache**
 - O(32) entries, each points to 256MB!
 - O(1) cycles to access
- **HW searches SLB**
 - To find ESID=>VSID mapping
 - Each STE (segment table entry) contains ESID & VSID info
- **OS explicitly manages SLB**
 - Maintains ESID=>VSID for all segments of the process
- **Upon SLB miss**
 - HW raises “segment fault” interrupt
 - OS will then insert right STE
- **Upon context switch**
 - OS invalidates SLB (not shared)

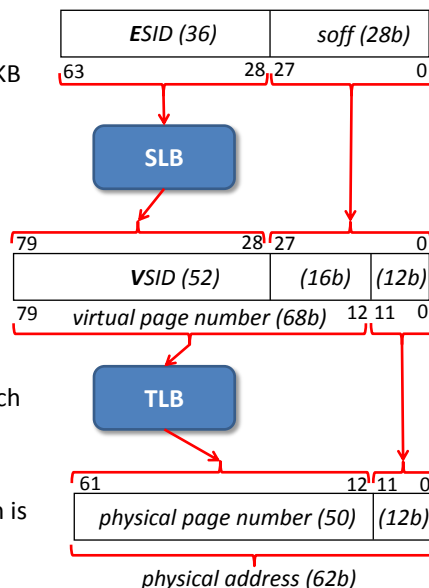


OS (234123) - virtual memory

47

PPC TLB (translation lookaside buffer)

- **TLB is a (less) fast HW cache**
 - O(1024) entries, each points to 4KB
 - O(10) cycles to access
- **HW searches TLB**
 - To find VPN=>PPN mapping
 - Each PTE (page table entry) contains VPN & PPN info
- **Shared by all processes**
 - Since virtual space is shared
 - Unlike in x86
 - Not invalidated upon context switch
- **Managed by HW & OS**
 - Upon miss, HW populates TLB
 - By “walking the page table”, which is in fact the “HTAB”...



OS (234123) - virtual memory

48

PPC HTAB (hash table)

- **PPC PTE (page table entry)**
 - Contains VPN (the “tag”; why do we need it?) & PPN
 - Unlike x86 PTE that only contains PPN
 - $|PTE| = 16 \text{ bytes} = 128 \text{ bit} > |VPN (68\text{bit})| + |PPN (50\text{bit})| = 118\text{bit}$
- **PTEG (page table entry group)**
 - Contains 8 PTEs
 - $8 * 16 = 128 \text{ bytes} = |\text{cache line}|$
- **HTAB**
 - At boot time, OS allocates in DRAM the “HTAB” array
 - Holds k PTEGs, where k is configurable
 - With a DRAM size of a handful of GBs:
 - It is recommended to set $|HTAB| = O(10\text{MB} - 100\text{MB})$
 - $\Rightarrow k = O(100,000 - 1,000,000)$
 - \Rightarrow HTAB points to $O(800,000 - 8,000,000)$ 4KB pages
 - \Rightarrow HTAB provides coverage of 3GB – 30GB of DRAM
 - OS saves HTAB’s size & base in “SDR1” (storage description register)

OS (234123) - virtual memory

49

PPC HTAB (hash table)

- **Upon TLB miss**
 - HW hashes VPN (modulo k)
 - Recall that HW knows about HTAB location and size via SDR1
 - The hash function is well-known and documented in the PPC spec
 - HW accesses HTAB and gets the so called “primary” PTEG of the VPN
 - HW searches for the VPN in the 8 PTEs populating the primary PTEG
 - If found, HW puts $VPN \Rightarrow PPN$ in TLB and re-executes operation
 - Otherwise, HW uses a secondary well-known hash function to obtain the “secondary” PTEG
 - If VPN found, HW puts $VPN \Rightarrow PPN$ in TLB and re-executes operation
 - Otherwise, HW triggers page-fault interrupt
 - OS will then resolve the page fault and put the appropriate PTE in one of the associated two PTEGs (primary or secondary)
 - After interrupt is handled, HW will re-execute the operation; now it’ll find the VPN in one of the 2 PTEGs

OS (234123) - virtual memory

50

PPC HTAB (hash table)

- **HW searches for VPN=>PPN translation as follows:**
 - if(VPN found in TLB)
 - Get PPN from TLB
 - else if(VPN found in primary_PTEG = HTAB[hash1(VPN) % k])
 - Get PPN from primary_PTEG
 - else if(VPN found in secondary_PTEG = HTAB[hash2(VPN) % k])
 - Get PPN from secondary_PTEG
 - else
 - Generate page-fault interrupt
 - OS will place appropriate PTE in either primary or secondary PTEGs
 - Next time, faulting operation will succeed reloading

OS (234123) - virtual memory

51

PPC ERAT (effective to real translation)

- **A small, fast HW cache**
 - O(128) entries
 - Quicker than 1-cycle to access (on the critical path)
 - Translates from effective to real (physical)
 - Analogous to x86 TLB (L1)
 - Updated to hold the most recent effective=>physical mappings used
 - On a LRU basis
 - If hit, don't need to go through the SLB/TLB process
 - In most cases, obviates the need to do costly SLB=>TLB=>HTAB translations

OS (234123) - virtual memory

52


Actually, PPC uses TLB hierarchy

- **ERAT**
 - TLB L1
- **TLB**
 - TLB L2
- **HTAB**
 - TLB L3
- **Only if not found in all 3 levels**
 - Go to OS, which has all information
- **Likewise, Intel & AMD**
 - Introduced a (bigger, slower) TLB 2
- **The reason...**
 - DRAM is getting bigger
 - TLB L1 remains roughly the same size

OS (234123) - virtual memory

53

TLB coverage drops exponentially



<http://static.usenix.org/event/osdi02/tech/navarro.html>

OS (234123) - virtual memory

54

How to increase TLB converge?

- **Superpages (also called huge pages): page size > 4KB**
 - Different sizes supported
 - Specifics vary among different architectures
 - From a few MBs to a few GBs!
(E.g., Intel supports 2MB, 4MB, 1GB)
- **TLB hierarchy**
 - Like the caches L1, L2, ...
 - Architectures can also support TLB-L1 (size = a few dozens), TLB-L2 (size = a few hundreds)
 - TLB-L2 is bigger, but is slower